

Generalized Alternation in Shared-Memory Systems

Michael E. Goldsby *

December 15, 2021

Abstract

This paper presents a new algorithm for generalized alternation in shared-memory systems. The algorithm is fully decentralized and contains no timed waits. Included are a proof sketch and a CSPM (machine-readable CSP) model of the algorithm suitable for use with the FDR refinement checker.

1 Introduction

1.1 Alternation

Hoare introduced the alternation construct in his first paper on communicating sequential processes (CSP) [Hoa78]. The alternation is a generalization of the `if...then...else` statement that allows the synchronization of processes and communication among them. It is based on Dijkstra's guarded commands [Dij75] but differs by allowing synchronization and communication as part of the guard.

Under development by Hoare and others, CSP went on to become a full-fledged process algebra [BHR84]. Here we follow the CSP scheme of representing a concurrent computation as a collection of processes that communicate by means of message-passing over synchronous channels. Each process

*michaelegoldsby@gmail.com

contains local variables that it alone may alter, and any global data must be invariant.

An alternation (or *alt* for short), like an `if...then...else` statement, consists of one or more branches. Each branch contains a *guard* and a *body*. The guard contains an optional boolean expression (the *condition*) and a *synchronization*. The condition may be absent, equivalent to its being true. The synchronization may be an *input*, an *output* or a *skip* (no synchronization).

The CSP algebra allows multi-way events and allows two-way data passage over the channels, but languages inspired by it (`occam`, `go`) typically allow only one-way communication between a sender and a receiver over a one-way synchronous channel.¹ Here also we follow the convention that a channel is a one-way point-to-point connection between a pair of processes and synchronization must involve one input (represented by `?`) and one output (represented by `!`). The communication of data may be absent, making the interaction a pure synchronization.

The body contains executable statements and is the only part of the branch that may change the local variables of the process. It consists of assignment statements and procedure calls, and we may also refer to it as “the assignments”. Note that the body is the only part of an *alt* that may change the values of the process’s local variables.

An alternation completes when exactly one of its branches completes (and transfers the data if any). If more than one branch is capable of completion, the choice is nondeterministic.

Implementation of alternation is simpler if the guards may contain only inputs (or only outputs), that is, if one side of the synchronization/communication is committed. An alternation construct that allows either input or output in its guards is said to be generalized. Generalized alternation removes the non-essential asymmetry of input-only alternation and allows a simpler expression of certain algorithms.

In addition to channel guards and skip guards, a branch may have a time-out guard. For the sake of simplicity, we restrict ourselves to channel guards in this paper. Also for simplicity’s sake, we assume that processes are non-terminating. For similar reasons we omit the details of data communication, as explained in section 2 below.

Figure 1 shows a symbolic representation of an alternation with two branches.

¹An exception is `occam-pi`, with its barrier synchronization [BWS05].

```

(
  condition1 ; c1?var1 -> ..body1..
[] condition2 ; c2!expr2 -> ..body2..
)

```

Figure 1: An alternation with two branches.

1.2 Background

Alternation is a useful construct even outside the context of CSP. The current work grew out of projects that required writing software for embedded microcontrollers. The initial work used a single processor, serving processes in round-robin fashion and running their branches to completion without interruption, which imposed some delay on interrupt handling. In that setting, it was simple to implement generalized alternation.

When extending the work to multicore systems (and systems with more immediate interrupt handling), the author did not wish to give up the symmetry and convenience provided by generalized alternation, which led to the work presented here.

Section 2 describes the algorithm and section 3 presents a proof sketch for it. Section 4 comments briefly on performance, section 5 lists related work, and section 6 describes possible future work. The appendices contain a CSPM model that the author used with the FDR4 [Gib+14; Gib21; Gib+13] refinement checker to guide development of the algorithm.

2 Description

Figure 2 shows pseudo-code for the algorithm.

```

1
2
3   ProcState = { Alting , Completing , Quiescent , Clearing }
4
5   Process = {
6       state: ProcState = Alting
7       alt: Alternation

```

```

8           mutex: Mutex
9       }
10
11   Channel = {
12       mutex: Mutex
13       first , second: Process
14   }
15
16   Alternation = {
17       nb: int
18       [nb] branch: Branch
19   }
20
21   Branch = {
22       c: Channel
23       partner: Process
24   }
25
26
27   run_branch(self: Process , partner: Process , b; int) :
28       (fired: bool) =
29
30       IF self.state = Alting THEN
31           IF c.first = nil and c.second = nil THEN
32               .. first
33               c.first , fired := self , false
34
35           ELSE IF c.first  $\diamond$  nil and c.first  $\diamond$  self THEN
36               IF partner.state = Alting THEN
37                   ..become second and complete
38                   c.second , self.state , partner.state :=
39                       self Completing , Quiescent
40                   ..do xfr if any
41                   fired := true
42                   ..run assignments
43               ELSE
44                   ..this alt completed elsewhere
45                   self.state , fired := Clearing , false

```

```

46             END
47         ELSE
48             ..ignore anything else
49             fired := false
50         END
51     END
52
53     ELSE IF state = Quiescent THEN
54         ..ignore everything
55         fired := false
56     END
57
58     ELSE IF state = Completing THEN
59         IF c.first = self and second <> nil THEN
60             .. first completes
61             c.first , c.secod := nil , nil
62             self.state , partner , fired :=
63                 Clearing , Clearing , true
64             ..run assignments
65         END
66     END
67 END
68
69 ELSE ..state = Clearing
70     IF c.first = self and c.second = nil THEN
71         c.first , fired := nil , false
72
73     ELSE ..ignore anything else
74         fired := false
75     END
76 END
77
78 run_proc( self: Process ) =
79     c := self.alt.c[b]
80     partner := self.alt[b].partner
81
82     proc1 , proc2 := IF self < partner
83                     THEN self , partner

```

```

84                                     ELSE partner , self END
85
86     nb := self.alt.nb
87     count: int = 0
88     WHILE true
89         b: int = 0
90         WHILE b < nb
91             c.mutex.lock
92             proc1.mutex.lock
93             proc2.mutex.lock
94             fired := run_branch(self , partner , b)
95             proc2.mutex.unlock
96             proc1.mutex.unlock
97             c.mutex.unlock
98             b, count := b+1,
99                 IF not fired THEN count+1 ELSE 0 END
100         END
101         ..b = nb
102         IF self.state = Clearing THEN
103             self.state := Alting END
104     END

```

Figure 2. Pseudo-code for the generalized alternation algorithm.

We regard a stand-alone input or output to be a degenerate case of alternation, so all interprocess communication and synchronization is performed by alternation.² Thus the execution of a process involves the successive execution of a series of alts, possibly interspersed by non-alt, non-communication/non-synchronizaton code..

Also assume for clarity that each process runs on its own processor, whose scheduler calls the `run_proc` method, although nothing in the algorithm need change if several processes were (fairly [Fra86]) scheduled on the same processor.

We can remove the boolean conditions from consideration. The body of a branch can alter the values of the variables on which a guard depends,

²Interestingly, any CSP process has an implementation as a single iterated alt [ABC87]. We make no use of that fact in this paper.

changing its truth value. However, since by our assumption a process can install a different alt when it finishes its previous one, changing the values of the guards is equivalent to installing an alt lacking the branches with falsified conditions.

Synchronization over a channel must involve exactly two processes, a sending process and a receiving process. We term the two processes *complementary* and say that one of the process is the *complementary process* or the *complement* of the other with respect to the channel. We may also simply refer to the complementary process when the channel is understood. We refer to the two complementary processes as *partners* wrt to the channel between them. If the alts of two processes share a channel, we call them *related* processes.

We say a process *selects* a channel if it returns **true** from the **run_branch** procedure when the process calls it for a branch containing the channel. A **true** return from **run_branch** means the branch will execute its body, possibly changing the values of the process's local variables.

The first and second fields of all channels are initially nil, and initially the state of all processes is **Alting**.

Suppose processes P and Q synchronize on a channel c. The sequence of events for such a synchronization is as follows:

- One of the processes, say P, finds $c.first = nil$. We say that P is first to the channel. P sets $c.first = P$ and carries on executing its branches (see lines 30-33 of Fig. 2).
- P's partner Q discovers that $c.first$ is not nil and is not itself (so it must be Q's partner wrt c) and also that $c.second$ is nil. Provided P's state is **Alting**, it then stores itself in $c.second$ and returns **true**, selecting c. We say that Q is second to the channel and has made an unmatched selection of c. (See lines 35-42 of Fig. 2.) If P's state is not **Alting**, the alt must have completed in another branch, and P sets its state to **Clearing** (see lines 43-45).

If the branch includes data transfer as well as synchronization, Q carries it out before returning **true**. Both sender and receiver have at this point visited c and can have left sufficient information in c to allow the transfer to occur (source address, destination address and data length.) For simplicity's sake we do not show these details here.

Note that it does not matter which of P and Q is the sender and which is the receiver; except for the actual transfer, the algorithm is symmetric wrt to sender and receiver.

When it returns **true**, we say that Q has completed as second to the channel. Note that the second process to the channel completes first.

Before returning, Q sets its own state to **Quiescent** and P's state to **Completing**.

- Next, P discovers that `c.first = itself` and `c.second` is non-nil. In response, it will set both fields to nil and return **true**, selecting `c`. We say now that P completes as first, and that the selection matches the earlier (unmatched) selection of its partner.³

Before returning, P sets both its own and its partner Q's state to **Clearing**. (See lines 58-66 of Fig. 2)

- In the **Clearing** state, a process visits each branch of its alt (by calling `run_branch`) and if it finds that it is first in the branch's channel, it resets the first field to nil (lines 69-71 of Fig. 2).

3 Proof sketch

The correctness criteria consist of a liveness property and a safety property.⁴

Note that during the execution of a branch, the executing process has exclusive access to itself, its partner and the channel the partners share (lines 91-97 of Fig. 2). Thus it is assured that the process's partner is not executing its `run_branch` procedure. Also note that we always lock the partners' mutexes in the same order to avoid an AB-BA deadlock [Han73] (lines 82-84 of Fig. 2).

³Note that an arbitrary interval of time may separate the second process's and the first process's completions, even though these two events together constitute the synchronization of the two processes. In practice, this interval would tend to be short.

⁴Any correctness condition is expressible as the conjunction of a liveness condition and a safety condition [Sch97].

3.1 Liveness

L1. If the system contains any complementary pair both of whose partners are in an alt, some alt will eventually complete; that is, for some complementary pair R, S , R will complete as second and S will complete as first.

Let P and Q be partners wrt channel c , and let c appear in both of their alts.

L0.1 If $c.\text{first} = \text{nil}$, then one of P, Q will become first to c unless P, Q or one of their partners selects some other channel.

L0.2 If $c.\text{first}$ is neither P nor nil and $c.\text{second} = \text{nil}$, then P will complete as second to c unless P or Q completes on some other channel.

L0.3 If $c.\text{first} = P$ and $c.\text{second}$ is not nil , then P will eventually select c and complete as first to c . When it does, it resets $c.\text{first}$ and $c.\text{second}$ to nil values.

Since each process continually polls its branches (lines 88-100 of Figure 2), L0.1-L0.3 follow directly from the code, and L1 follows from L0.1-L0.3.

3.2 Safety

S1. If a process P completes as second wrt c , the next selection, if any, produced by its complementary process Q wrt c will be of c .

S2. Until Q performs that selection, P will perform no other selection.

That is, selections occur in matched pairs. ⁵

3.3 Proof sketch for safety conditions

S1: When P completes as second, it puts its complementary process Q wrt c into the **Completing** state; it leaves that state only when it completes as first to c .

⁵But note that completions of unrelated processes may be interspersed with each other.

S2: When P completes as second, it enters the **Quiescent** state, which it leaves only when Q puts it into the **Clearing** state, which it does only when Q completes as first.

The above considerations, coupled with the observation that a process can only complete (as first or second) when in the **Alting** state, imply S1 and S2.

4 Performance

Assume that processes P and Q complete their alts for channel c , that P's alt has m branches and Q's has n , and that P is first to the channel. P must execute at most m branches to become first, after which Q must execute at most n branches to become second and complete, then again P must execute at most m branches to complete as first. Then process P must execute m branches and Q must execute n branches in order to clear their channels. Thus the algorithm executes at most $3m + 2n$ branches to complete the alt. Note, however, that many of the branch executions of P and Q may occur concurrently. It is only when a branch contains a channel c' such that P and Q are partners wrt c' that P and Q cannot proceed concurrently.

The execution of the body of a branch may take any amount of time, but it takes place outside any locks and can therefore proceed concurrently with its partner's execution (and that of every other process).

The biggest performance concern is the three locks, for c , P, and Q, that the program must acquire in order to run a branch. The author looks forward to making performance studies with an actual implementation (see Section 6.)

5 Related work

Most generalized alternation algorithms in the literature are for systems in which the processes communicate via message passing [BK84; Bag86; Ber80; BHR84; BS83; Dem98; BK92; Sch82; Sil79]. Fujimoto and Feng [FF87] present an algorithm for shared-memory systems with a proof of its correctness. It contains a timed delay and uses interprocessor signals in addition to locks.

It might be possible to use a message-passing algorithm as a model for a shared-memory algorithm by substituting procedure calls for messages. However, the patterns of mutual exclusion required to make the procedure calls atomic would tend to be costly, and most of the cited algorithms have other drawbacks, such as high message or time complexity, or a tendency to deadlock under certain messages patterns. The best candidate for such a conversion appears to be an efficient algorithm in [Dem98].

6 Future work

Presently, the partner processes synchronize at branch resolution. We might try to make the span of the synchronization smaller, possibly by use of the compare-and-swap instruction.

We could perhaps improve the algorithm's power usage by requiring it, after it has visited every branch of its alt, to wait for a signal from a partner before executing its branches again.

We might also try to lessen the number of branches that the first process to the channel must execute in order to complete by letting the second process specify the branch that the first is to execute, in effect “forcing” a branch on it.

We may undertake a formal proof of the algorithm.

We have the intention to make a C or C++ implementation of the algorithm and study its performance.

We might investigate the possibility of using this shared-memory algorithm as a pattern for the development of a message-passing algorithm, although at a message per procedure call, the message complexity does not appear favorable.

References

- [ABC87] K. R. Apt, Luc Bougé, and Ph. Clermont. “Two normal form theorems for CSP programs”. In: *Information Processing Letters* 26.4 (1987), pp. 165–171.

- [Bag86] R. Bagrodia. “A distributed algorithm to implement the generalized alternation command in CSP”. In: *The 6th International Conference on Distributed Computing Systems*. May 1986, pp. 422–427.
- [Ber80] A. J. Bernstein. “Output guards and nondeterminism in communicating sequential processes”. In: *ToPLaS* 2.2 (Apr. 1980), pp. 234–238.
- [BHR84] Stephen Brookes, C. A. R. Hoare, and A. W. Roscoe. “A theory of communicating sequential processes”. In: *JACM* 31.3 (July 1984), pp. 560–599.
- [BK84] Eklund Back and Kurki-Suonio. *A fair and efficient implementation of CSP with output guards*. Tech. rep. A-38. Department of Computer Science, University of Utah, 1984.
- [BK92] Eklund Back and Kurki-Suonio. *A distributed protocol for channel-based communication with choice*. Tech. rep. ECRC-92-16. Muenchen: European computer-industry research center, 1992.
- [BS83] G. N. Buckley and A. Silberschatz. “An effective implementation for the generalized input–output construct of CSP”. In: *ToPLaS* 5.2 (Apr. 1983), pp. 223–235.
- [BWS05] Frederick R. M. Barnes, Peter H. Welch, and Adam T. Sampson. “Barrier synchronization for occam-pi”. In: *Proceedings of the International Conference on Parallel and Distributed Processing Techniques and Applications, PDPTA 2005, Volume 1*. Ed. by Hamid R. Arabnia. Las Vegas, NV, June 2005, pp. 173–179.
- [Dem98] Erik D. Demaine. “Protocols for non-deterministic communication over synchronous channels”. In: *12th International Parallel Processing Symposium / 9th Symposium on Parallel and Distributed Processing (IPPS/SPDP '98)*. Orlando, FL: IEEE Computer Society, 1998, pp. 24–30.
- [Dij75] E. W. Dijkstra. “Guarded commands, nondeterminism and formal derivation of programs”. In: *CACM* 18.8 (Aug. 1975), pp. 453–457.

- [FF87] Richard M. Fujimoto and Hwa-chung Feng. *A shared memory algorithm and proof for the alternative construct in CSP*. Tech. rep. UUCS-87-021. Department of Computer Science, University of Utah, 1987.
- [Fra86] Nissim Francez. *Fairness*. Springer Verlag, 1986.
- [Gib+13] Thomas Gibson-Robinson et al. *Failures Divergences Refinement (FDR) Version 3*. 2013. URL: <https://www.cs.ox.ac.uk/projects/fdr/>.
- [Gib+14] Thomas Gibson-Robinson et al. “FDR3 — A Modern Refinement Checker for CSP”. In: *Tools and Algorithms for the Construction and Analysis of Systems*. Ed. by Erika Ábrahám and Klaus Havelund. Vol. 8413. Lecture Notes in Computer Science. 2014, pp. 187–201.
- [Gib21] Thomas Gibson-Robinson. *FDR Manual*. 4.2.7. Accessed December, 2021, at <https://www.cocotec.io/fdr/fdr-manual.pdf>. University of Oxford. May 2021.
- [Han73] Per Brinch Hansen. *Operating System Principles*. Englewood, NJ: Prentice-Hall, 1973.
- [Hoa78] C. A. R. Hoare. “Communicating sequential processes”. In: *CACM* 21.8 (Aug. 1978), pp. 666–677.
- [Sch82] F. B. Schneider. “Synchronization in distributed programs”. In: *ToPLaS* 4.2 (Apr. 1982), pp. 125–148.
- [Sil79] A. Silberschatz. “Communication and synchronization in distributed systems”. In: *IEEE Transactions of Software Engineering* SE-5.6 (Nov. 1979), pp. 542–546.

Appendix A

The FDR4 refinement checker allowed the author to find errors when developing the algorithm. In the author’s judgement, it saved time and effort, though it is difficult to characterize how much, especially since it was necessary to climb a substantial learning curve in order to use the checker.

This appendix contains a CSPM model of the algorithm for use with FDR4. An included file contains functions that define the process network. The author exploits this separation in a somewhat crude but relatively effective way: the inclusion of a new network file in general causes a rash of trivial compile-time errors, which the author then corrects by hand. The specifications have to be rewritten with some care, however, for each new network configuration.

Since CSP has no notion of fairness, it is necessary in this model to include a scheduler (the SCHED process) to make sure that each CSP process gets a chance to execute.

```
1
2  --CHAN and PROC ids
3  NONE = 0
4
5  CHAN_C = 1
6  CHAN_D = 2
7  CHAN_E = 3
8
9  PROC_P = 1
10 PROC_Q = 2
11
12 --state
13 ALTING = 1
14 COMPLETING = 2
15 QUIESCENT = 3
16 CLEARING = 4
17
18
19 --CHAN and PROC ids
20 --channels bet. RUNBRANCH and CHAN
21 channel pc, qc, pd, qd, pe, qe, chanerr: ChanOp
```

```

22
23  —channels bet. RUN_PROC or RUN_BRANCH and own PROC
24  channel p-p, q-q, procerr: ProcOp
25
26  —channels bet. RUN_BRANCH and partner's PROC
27  channel p-q, q-p: PartnerOp
28
29  —channels bet. RUN_PROC and RUN_BRANCH:
30  channel p-run, q-run, runerr: RunOp
31
32  —channel between SCHED and RUN_PROCS
33  channel schp, schq, scherr: SchedOp
34
35  nametype ProcId = {0..5}
36  nametype ChanId = {0..4}
37  nametype BranchNo = {0..3}
38  nametype ErrNo = {0..7}
39  nametype ProcState =
40      { ALTING, COMPLETING, QUIESCENT, CLEARING }
41  nametype Position = {2..36}
42
43  datatype ChanOp = RD_1ST.ProcId | WR_1ST.ProcId |
44                  RD_2ND.ProcId | WR_2ND.ProcId |
45                  RD_BR.BranchNo | WR_BR.BranchNo |
46                  LOCK_CHAN | UNLOCK_CHAN |
47                  C_ERR.ErrNo
48  datatype ProcOp = WR_CURR_BR.BranchNo | RD_CURR_BR.BranchNo |
49                  RD_NB.BranchNo | RD_STATE.ProcState |
50                  WR_STATE.ProcState | WR_FIRED.Bool |
51                  RD_FIRED.Bool | LOCK_PROC.ProcId
52                  | UNLOCK_PROC | P_ERR.ErrNo
53  subtype PartnerOp = WR_CURR_BR.BranchNo | WR_STATE.ProcState |
54                  RD_STATE.ProcState | LOCK_PROC.ProcId |
55                  UNLOCK_PROC
56  datatype RunOp = CALL.BranchNo | RET.Bool.ChanId
57  datatype SchedOp = GO | DONE
58
59  —topology-dependent functions

```

```

60         include "2p3c.csp"
61
62     --runs a designated branch of the alternation
63     RUN_BRANCH(pid) =
64     (
65         -- local variables
66         -- pid proc id of self
67
68         let
69             self = pid2proc(pid)
70             run = pid2run(pid)
71         within
72             run.CALL?b
73
74         -> let
75             cid = br2cid(pid, b)
76             c = ids2chan(pid, cid)
77             partner = ids2partner(pid, cid)
78         within
79
80             self.RD.STATE?state
81             -> c.RD.1ST?first
82             -> c.RD.2ND?second
83
84             -> if (state==ALTING) then
85
86                 if ((first==NONE) and (second==NONE)) then
87                     --first
88                     c.WR.BR!b
89                     -> c.WR.1ST!pid
90                     -> run.RET!false!cid
91                     -> RUN_BRANCH(pid)
92
93                 else if ((first!=NONE) and (first!=pid)
94                     and (second==NONE)) then
95                     partner.RD.STATE?pstate
96                     -> if (pstate==ALTING) then
97                         --second completes

```



```

98         c.WR_2ND!pid
99         -> c.RD_BR?b1
100        -> partner.WR.STATE!COMPLETING
101        --(do data transfer if any)
102        -> self.WR.STATE!QUIESCENT
103        -> run.RET!true!cid
104        -> RUN_BRANCH(pid)
105    else
106        --alt completed elsewhere
107        self.WR.STATE!CLEARING
108        -> run.RET!false!cid
109        -> RUN_BRANCH(pid)
110
111    else --ignore anything else
112        --procerr.P.ERR!1
113        run.RET!false!cid
114        -> RUN_BRANCH(pid)
115
116    else if (state==QUIESCENT) then
117        run.RET!false!cid
118        -> RUN_BRANCH(pid)
119
120    else if (state==COMPLETING) then
121
122        if ((first==pid) and (second!=NONE)) then
123            --first completes
124            c.WR_1ST!NONE
125            -> c.WR_2ND!NONE
126            -> self.WR.STATE!CLEARING
127            -> run.RET!true!cid
128            -> partner.WR.STATE!CLEARING
129            --(run assignments)
130            -> RUN_BRANCH(pid)
131
132        else --ignore everytning else
133            run.RET!false!cid
134            -> RUN_BRANCH(pid)
135

```

```

136         else —state = Clearing
137             if ((first==pid) and (second==NONE)) then
138                 c.WR_1ST!NONE
139                 -> run.RET!false!cid
140                 -> RUN_BRANCH(pid)
141
142             —ignore any other combination
143             else
144                 run.RET!false!cid
145                 -> RUN_BRANCH(pid)
146
147     )
148
149
150
151 RUNPROC( pid, b, count, first, state ) =
152 (
153     — parameters:
154     —     pid          proc id
155     — local variables:
156     —     b            branch number in 0..nb-1
157     —     count        progress counter for branch processing
158     —     first        start of execution for a process
159     —     state        state (ALTING, COMPLETING or CLEARING)
160     if (first) then
161         let
162             sch = pid2sch(pid)
163         within
164             sch.GO
165             -> RUNPROC( pid, b, count, false, state )
166     else
167         let
168             self = pid2proc(pid)
169             cid = br2cid(pid,b)
170             c = ids2chan(pid,cid)
171             partid = ids2partid(pid,cid)
172             partner = ids2partner(pid,cid)
173             (proc1, proc2) = if pid < partid

```

```

174                                     then (self , partner)
175                                     else (partner , self)
176     nb = pid2nb(pid)
177     run = pid2run(pid)
178 within
179     c.LOCK_CHAN
180     -> proc1.LOCK_PROC.pid
181     -> proc2.LOCK_PROC.pid
182     -> if count < nb then
183         run.CALL!b
184         —between CALL and RET, RUN.BRANCH
185         —has the use of x_y, x = pid,
186         —y in { y | y a partner of x } U {x}
187         -> run.RET?fired?cid
188         -> proc2.UNLOCK_PROC
189         -> proc1.UNLOCK_PROC
190         -> c.UNLOCK_CHAN
191         -> if (fired) then
192             —clear all branches
193             RUN_PROC( pid,
194                 ((b+1)%nb), 0, first , state )
195         else
196             RUN_PROC( pid,
197                 ((b+1)%nb), count+1, first , state )
198     else —count = nb
199         — branches exhausted
200         self.RD.STATE?state
201         -> if state==CLEARING then
202             self.WR.STATE!ALTING
203             -> proc2.UNLOCK_PROC
204             -> proc1.UNLOCK_PROC
205             -> c.UNLOCK_CHAN
206             -> RUN_PROC( pid,
207                 b, 0, true, ALTING )
208         else —keep Alting
209             proc2.UNLOCK_PROC
210             -> proc1.UNLOCK_PROC
211             -> c.UNLOCK_CHAN

```

```

212                                     -> RUNPROC( pid,
213                                     b, 0, true, state )
214
215 )
216
217 PROC( pid, nb, b, fired, state, locked ) =
218 (
219     —parameters:
220     —   pid      id of this process
221     —   nb      # of branches in process's alt
222     —local variables:
223     —   b      branch number for completion
224     —   state   Alting or Clearing
225     —   fired   true if most recently-run branch fired
226     —   locked  true if locked
227
228     let
229         p = pid2proc(pid)
230         pp = pid2partners(pid)
231     within
232
233         (not locked) & p.LOCK_PROC?id
234         -> PROC( pid, nb, b, fired, state, true )
235     []
236         (locked) & p.UNLOCK_PROC
237         -> PROC( pid, nb, b, fired, state, false )
238     []
239         p.WR_CURR_BR?newb
240         -> PROC( pid, nb, newb, fired, state, locked )
241     []
242         p.RD_CURR_BR!b
243         -> PROC( pid, nb, b, fired, state, locked )
244     []
245         p.RD_NB!nb
246         -> PROC( pid, nb, b, fired, state, locked )
247     []
248         p.WR_STATE?newstate
249         -> PROC( pid, nb, b, fired, newstate, locked )

```

```

250     []
251     p.RD_STATE!state
252     -> PROC( pid, nb, b, fired, state, locked )
253     []
254     p.WR_FIRED?newfired
255     -> PROC( pid, nb, b, newfired, state, locked )
256     []
257     p.RD_FIRED!fired
258     -> PROC( pid, nb, b, fired, state, locked )
259     []
260     [] c: pp @ c.WR_CURR_BR?newb
261     -> PROC( pid, nb, newb, fired, state, locked )
262     []
263     [] c: pp @ (not locked) & c.LOCK_PROC?id
264     -> PROC( pid, nb, b, fired, state, true )
265     []
266     [] c: pp @ (locked) & c.UNLOCK_PROC
267     -> PROC( pid, nb, b, fired, state, false )
268     []
269     [] c: pp @ c.RD_STATE!state
270     -> PROC( pid, nb, b, fired, state, locked )
271     []
272     [] c: pp @ c.WR_STATE?newstate
273     -> PROC( pid, nb, b, fired, newstate, locked )
274
275
276 )
277
278 CHAN( cid, first, second, b, locked ) =
279 (
280     —parameters:
281     —   cid:          id of this channel
282     —               (CHAN_C or CHAN_D, eg)
283     —local variables:
284     —   first         first process to the channel
285     —   second        second ...
286     —   b             branch number for first
287     —   locked        chan locked when true

```

```

288 let
289     (c1,c2) = cid2procs(cid)
290 within
291
292     (not locked) & c1.LOCK_CHAN
293     -> CHAN( cid , first , second , b, true )
294 []
295     (locked) & c1.UNLOCK_CHAN
296     -> CHAN( cid , first , second , b, false )
297 []
298     c1.WR_1ST?newfirst
299     -> CHAN( cid , newfirst , second , b, locked )
300 []
301     c1.RD_1ST!first
302     -> CHAN( cid , first , second , b, locked )
303 []
304     c1.WR_2ND?newsecond
305     -> CHAN( cid , first , newsecond , b, locked )
306 []
307     c1.RD_2ND!second
308     -> CHAN( cid , first , second , b, locked )
309 []
310     c1.WR_BR?newb
311     -> CHAN( cid , first , second , newb, locked )
312 []
313     c1.RD_BR!b
314     -> CHAN( cid , first , second , b, locked )
315
316 []
317
318     (not locked) & c2.LOCK_CHAN
319     -> CHAN( cid , first , second , b, true )
320 []
321     (locked) & c2.UNLOCK_CHAN
322     -> CHAN( cid , first , second , b, false )
323 []
324     c2.WR_1ST?newfirst
325     -> CHAN( cid , newfirst , second , b, locked )

```

```

326     []
327         c2.RD_1ST!first
328         -> CHAN( cid , first , second , b, locked )
329     []
330         c2.WR_2ND?newsecond
331         -> CHAN( cid , first , newsecond , b, locked )
332     []
333         c2.RD_2ND!second
334         -> CHAN( cid , first , second , b, locked )
335     []
336         c2.WR_BR?newb
337         -> CHAN( cid , first , second , newb, locked )
338     []
339         c2.RD_BR!b
340         -> CHAN( cid , first , second , b, locked )
341
342 )
343     {--
344     --sequential scheduler:
345     --let each process execute in turn until
346     --it relinquishes
347     SCHED( sch ) =
348     (
349         --parameters:
350         -- sch          list of channels to RUN_PROC processes
351
352         let
353             ch = head(sch)
354         within
355             ch.GO
356             -> ch.DONE
357             -> SCHED( tail(sch)^<ch> )
358     )
359     --}
360
361 --let all process execute together
362 SCHED( schp , schq , pgo , qgo ) =
363 (

```

```

364         pgo & schp.GO -> SCHED( schp, schq, false, true )
365     []
366         qgo & schq.GO -> SCHED( schp, schq, true, false )
367
368 )
369 --start all RUNPROC processes and let each one execute
370 --according to the signals sent by the scheduler process
371
372 RUN() = ( ||| pid: all_pids() @
373           ( RUNPROC( pid, 0, 0, true, ALTING )
374             [| run_events() |]
375             RUN_BRANCH( pid ) ) ) )
376
377 PROCS() = ( ||| pid: all_pids() @
378            PROC( pid,
379                pid2nb(pid), 0, false, ALTING, false ) )
380
381 CHANS() = ( ||| cid: all_cids() @
382            CHAN( cid, NONE, NONE, 0, false ) )
383
384 RUN_ALL() = ( ( RUN() [| proc_events() |] PROCS() )
385              [| chan_events() |] CHANS()
386              )
387              [| sched_events() |]
388              SCHED( schp, schq, true, false )
389
390
391 SYSTEM = RUN_ALL() \ Hidden
392         _____
393
394
395 SPEC =
396     (
397     (
398         ( p_run.RET.True.CHAN_C -> q_run.RET.True.CHAN_C
399           -> SPEC )
400         |~|
401         ( q_run.RET.True.CHAN_C -> p_run.RET.True.CHAN_C

```



```

402             -> SPEC )
403         )
404         |~|
405         (
406             ( p_run.RET.True.CHAND -> q_run.RET.True.CHAND
407               -> SPEC )
408             |~| —CHAND (in either order)
409             ( q_run.RET.True.CHAND -> p_run.RET.True.CHAND
410               -> SPEC )
411         )
412         |~|
413         (
414             ( p_run.RET.True.CHANE -> q_run.RET.True.CHANE
415               -> SPEC )
416             |~| —CHANE (in either order)
417             ( q_run.RET.True.CHANE -> p_run.RET.True.CHANE
418               -> SPEC )
419         )
420     )
421
422     assert SYSTEM :[divergence-free]
423         :[partial order reduce]
424     assert RUN_ALL() :[deadlock free [F]]
425         :[partial order reduce]
426     assert SPEC [T= SYSTEM
427         :[partial order reduce]

```

Appendix B

Below are the topology-defining functions for a simple network with two processes that share three channels.

File 2p3c.csp

```

1
2
3  {-          C
4      /      \
5     /        \
6    /          \
7   [P]---D---[Q]
8    \          /
9     \        /
10    \         /
11     \       /
12  -}
13
14
15 —topology-dependent functions
16 —
17 — pid, cid -> channel to CHAN
18 ids2chan(pid, cid) =
19     if ((pid==PROC.P) and (cid==CHAN.C)) then pc else
20     if ((pid==PROC.P) and (cid==CHAN.D)) then pd else
21     if ((pid==PROC.P) and (cid==CHAN.E)) then pe else
22     if ((pid==PROC.Q) and (cid==CHAN.C)) then qc else
23     if ((pid==PROC.Q) and (cid==CHAN.D)) then qd else
24     if ((pid==PROC.Q) and (cid==CHAN.E)) then qe else
25     chanerr
26
27 — pid -> channel between PROC and RUN.BRANCH
28 pid2proc(pid) = if (pid==PROC.P) then p_p else
29                  if (pid==PROC.Q) then q_q else
30                  procerr
31
32 —pid, br -> cid (id of channel

```

```

33 — associated with a process's branch)
34 br2cid(pid, b) =
35     if ((pid==PROC_P) and (b==0)) then CHAN_C else
36     if ((pid==PROC_P) and (b==1)) then CHAN_D else
37     if ((pid==PROC_P) and (b==2)) then CHAN_E else
38     if ((pid==PROC_Q) and (b==0)) then CHAN_C else
39     if ((pid==PROC_Q) and (b==1)) then CHAN_D else
40     if ((pid==PROC_Q) and (b==2)) then CHAN_E else
41     NONE
42
43 —cid to (channel,channel) = channels bet. CHAN and its PROCs
44 cid2procs(cid) = if (cid==CHAN_C) then (pc,qc) else
45                  if (cid==CHAN_D) then (pd,qd) else
46                  if (cid==CHAN_E) then (pe,qe) else
47                  (chanerr,chanerr)
48
49 —pid -> set of channels to partners' PROCs
50 pid2partners(pid) = if (pid==PROC_P) then { q-p } else
51                    if (pid==PROC_Q) then { p-q } else
52                    { }
53
54 —pid, cid -> channel between RUN_BRANCH and partner's PROC
55 ids2partner(pid, cid) =
56     if ((pid==PROC_P) and (cid==CHAN_C)) then p-q else
57     if ((pid==PROC_P) and (cid==CHAN_D)) then p-q else
58     if ((pid==PROC_P) and (cid==CHAN_E)) then p-q else
59     if ((pid==PROC_Q) and (cid==CHAN_C)) then q-p else
60     if ((pid==PROC_Q) and (cid==CHAN_D)) then q-p else
61     if ((pid==PROC_Q) and (cid==CHAN_E)) then q-p else
62     procerr
63
64 —pid, cid -> id of partner
65 ids2partid(pid, cid) =
66     if ((pid==PROC_P) and (cid==CHAN_C)) then PROC_Q else
67     if ((pid==PROC_P) and (cid==CHAN_D)) then PROC_Q else
68     if ((pid==PROC_P) and (cid==CHAN_E)) then PROC_Q else
69     if ((pid==PROC_Q) and (cid==CHAN_C)) then PROC_P else
70     if ((pid==PROC_Q) and (cid==CHAN_D)) then PROC_P else

```

```

71         if ((pid==PROC_Q) and (cid==CHAN_E)) then PROC_P else
72         NONE
73
74 —pid -> nr of branches
75 pid2nb(pid) =
76         if (pid==PROC_P) then 3 else
77         if (pid==PROC_Q) then 3 else
78         0
79
80 pid2run(pid) = if (pid==PROC_P) then p_run else
81                 if (pid==PROC_Q) then q_run else
82                 runerr
83
84
85 pid2sch(pid) = if (pid==PROC_P) then schp else
86                 if (pid==PROC_Q) then schq else
87                 scherr
88
89 all_pids() = { PROC_P, PROC_Q }
90
91 all_cids() = { CHAN_C, CHAN_D, CHAN_E }
92
93 chan_events() = { | pc, qc, pd, qd, pe, qe | }
94
95 proc_events() = { | p-p, q-q, p-q, q-p | }
96
97 run_events() = { | p-run, q-run | }
98
99 sched_list() = < schp, schq >
100
101 sched_events() = { | schp, schq | }
102
103 ———
104
105 Events = { |
106             pc, qc, pd, qd, pe, qe,
107             p-p, q-q, p-q, q-p, p-run, q-run,
108             schp, schq

```

```
109         |}  
110  
111 Hidden = diff( union(Events, {| report |} ),  
112                 {| p_run.RET.true, q_run.RET.true |} )  
113
```